

Assignment: CNN and MNIST

David Bertoldi – 735213
email: d.bertoldi@campus.unimib.it

Department of Informatics, Systems and Communication

University of Milano-Bicocca



1 Inspecting the data

The MNIST dataset contains 70 000 images of handwritten digits (0 to 9) that have been size-normalized and centered in a square grid of pixels. Each image is a 28×28 array of floating-point numbers representing grayscale intensities ranging from 0 (black) to 255 (white).

The labels consist of a vector of values, corresponding to the digit classification categories 0 through 9.

The dataset is already divided into training and test sets, respectively with 60 000 and 10 000 samples.

Figure 1 shows an example of the population.




Figure 1: The first 10 samples of the train dataset

The training population presents a distribution with mean $\mu = 6\,000$ and standard deviation $\sigma \simeq 340$ and thus we didn't notice any important unbalance in the data. For this reason we assumed the data followed a distribution $X \sim U(\mu, \sigma)$ and no data augmentation on less populated classes was taken into account. Figure 2 shows the data distribution for both training and test datasets.

2 Preparing the data

Before training a FFNN using this images, encoded in 28×28 matrices with values from 0 to 255, we flattened them in arrays 1×784 and rescaled each value in the continuous interval $[0, 1]$. This encoding will be used in every section of this work: a flat array better suits the input layer of a FFNN and small values increases the efficiency in the calculations.

2.1 Data split

As noted in section 1, the dataset is divided into training and test samples. A validation subset is missing and thus




Figure 2: Histogram of the frequency of samples in the dataset

is retrieved from the training set: 15% of the images are randomly used for validation instead of training (along with their labels) for a total of 9 000 samples.

About labels, we encoded them in one-hot vectors so that the 1s are set in the index representing the numerical class.

3 Building the network and training

The aim of this section is to describe a CNN with less than 10 000 parameters that is able to classify with high level of accuracy the numbers from the dataset in 10 epochs with batches of 128 samples. Most of the choices were done according to the results coming from the training and prediction phases.

3.1 The network

The CNN presents a typical architecture formed by convolutional layers followed by pooling layers and ending with dense layers.

In particular there are 2 convolutional layers covering the whole 28×28 matrix, formed by $8 \times 3 \times 3$ filters, for a total dimension of $28 \times 28 \times 8$. These two layers are followed by a max pooling layer that halves the the width and height of the outcoming activation map. For this problem we tested both *max pooling* and *average pooling*; the first one performed slightly better (+0.2% in test accuracy): usually *average pooling* smooths out the image and the sharp features may be identified with more difficulty, while *max pooling* chooses the

white pixels of the image (in case of MNIST dataset, the pixels defining the handwritten digit). Although we noticed a slightly improvement using *max pooling*, the images are too small to actually benefit from the methods' differences. The structure continues with another one convolutional layer aligned with the 2D spatiality of the last pooling layer but doubled in the depth, that is $7 \times 7 \times 16$. The convolutional layer is reduced in spatiality by another *max pooling layer* $7 \times 7 \times 16$ and flattened in a 1D array of 784. The input flows to an output layer activated by *Softmax* function. Figure 3 summarizes the entire architecture and Table 1 highlights the number of parameters in each layer.

Layer	Size	Parameters
input	$28 \times 28 \times 1$	0
Conv2D-1	$28 \times 28 \times 8$	$(3 \cdot 3 \cdot 1 + 1) \cdot 8 = 80$
Conv2D-2	$28 \times 28 \times 8$	$(3 \cdot 3 \cdot 8 + 1) \cdot 8 = 584$
MaxPool-1	$14 \times 14 \times 8$	0
Conv2D-3	$14 \times 14 \times 16$	$(3 \cdot 3 \cdot 8 + 1) \cdot 16 = 1\,168$
MaxPool-2	$7 \times 7 \times 16$	0
Flatten	$1 \times 1 \times 784$	0
Dense	$1 \times 1 \times 10$	$(784 + 1) \cdot 10 = 7\,850$
Total		9 682

Table 1: Summary of the layers' dimensions and count of the parameters for each layer




Figure 3: Architecture of the CNN

About the convolutional layers, Tensorflow allows to specify the padding: no padding and padding with zeroes. The first one would have reduced the spatiality of each layer by 1, the second one preserves the dimensions by putting evenly 0s on the margins. Adding 0s with this particular dataset is not an issue, because most of the images (if not all of them) do not contain information along the margins. We benchmarked the performances of both methods and we noticed that the zero padding performed had a higher validation accuracy of +0.8%.

3.2 Training

The choice of the optimizer was among *RMSProp* and *Adam*. None of them have shown signs of getting stuck in local minimum regions but *Adam* performed better overall, with +0.9% on test accuracy.

The only regularization used is the *dropout* technique applied to the flattened layer with a rate of 20%. This allowed the model to not overfit too much and generalize better the problem. The choice of the probability to remove a neuron is based on the chart in Figure 4 where we moved the rate from 0% to 90% and logged the results for training, validation and test accuracy.




Figure 4: Loss

Even if the test accuracy reached its maximum with a 10% rate (98.8%), we chose a 20% drop rate in order to make the model underfit a little more without compromising too much the overall accuracy.

The behaviour of the model during the training phase is described in the plots of the loss and categorical accuracy in Figure 5 and Figure 6.




Figure 5: Loss

The validation loss was always lower than the training loss. This effect is caused by the *dropout* because it penal-




Figure 6: Categorical accuracy

izes model variance by randomly removing neurons from the flatten layer only during the training; the models underfitted and since *dropout* is disabled during the validation we had lower validation loss. The same for the accuracy, being higher during validation.

The best validation accuracy reached was 98.98% using *Adam* with learning rate of 10^{-3} and batch size of 256.

3.3 Evaluation

We tested the quality of the model's predictions using the provided test dataset, applying very same transformations on the images used for the training. The categorical accuracy over the test set reached 98.65% and can be analyzed with the help of the confusion matrix (Figure 7).




Figure 7: Confusion matrix over the test data

We noticed that the model confused the 2.34% of the 4s with a 9: this can be explained by the fact that the two digits have similar forms, in particular when the 4 is written quickly and in its "closed form". Another source of confusion is the 8 being confused again with a 9 (1% of the cases). Apart from

these cases, that have a relatively low impact, the model gave really good results.